

An (Interactive) Introduction to



SecFog

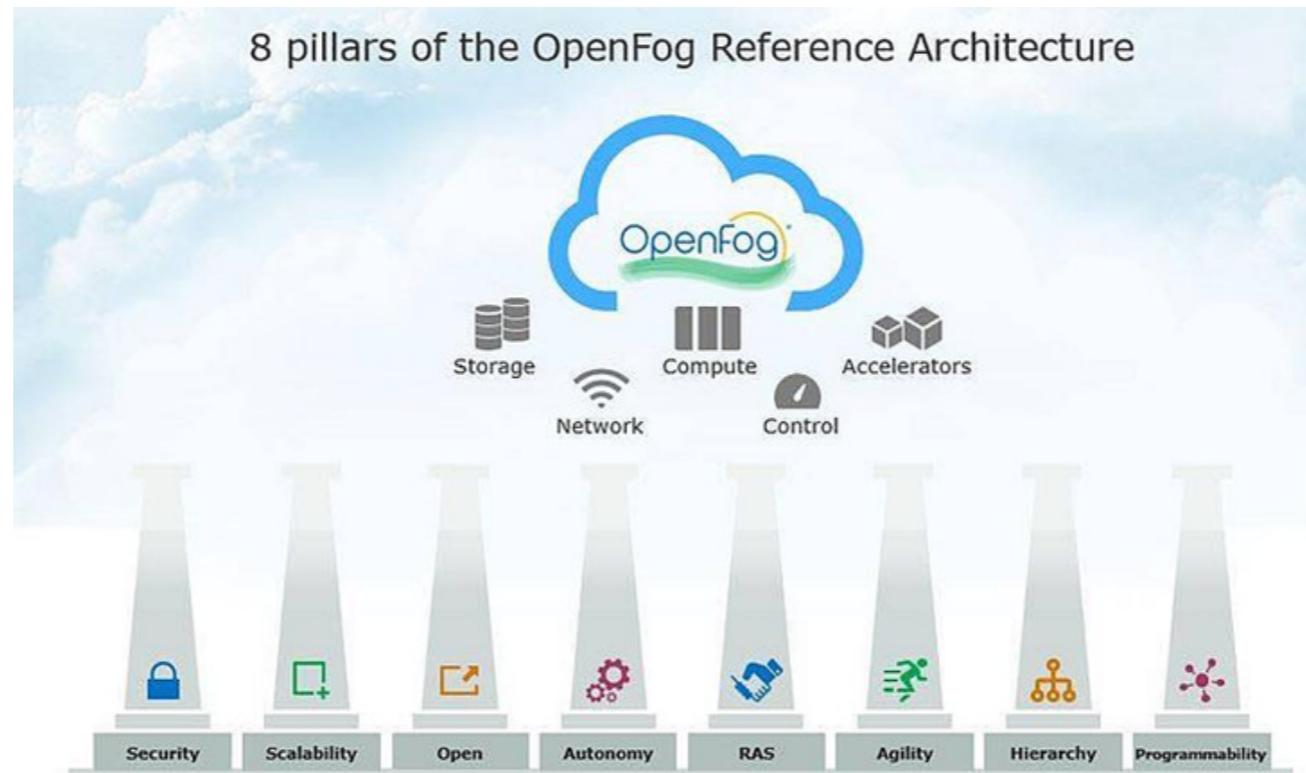
Stefano Forti

Department of Computer Science, University of Pisa, Italy

Reading: A. Brogi, S. Forti, G.-L. Ferrari, [Secure Apps in the Fog: Anything to Declare?](#), 2018.

Security in the Fog

- The Fog will **increase** the number of **security enforcement points**.
- The Fog will **share** with the Cloud many **security threats** (e.g., virtualisation-related).
- The Fog will be **exposed** to new, peculiar **threats** due to physical accessibility of Fog nodes.



Security is one of the pillars of OpenFog Reference Architecture.

Considered Problem

The problem is known as *Component Deployment Problem* (CDP).

Given

- a multi-component application A with requirements R
- a distributed Fog infrastructure I (nodes and links)
- a set of objective metrics M

determine eligible application deployments that meet all requirements in R and optimise metrics in M .



How difficult is CDP?

To prove P is NP-hard, we take another problem P' that is known to be NP-hard and we reduce P' to P in polynomial time, i.e. $P' \rightarrow_p P$.

This proves that P is at least as hard as P' . Hence, finding a poly-time algorithm that solves P would lead to solve P' in poly-time and prove that $P = NP$. (You can give it a try 😎)

Consider the **Subgraph Isomorphism Problem (SIP)** as our P' :

Given two graphs G and H , a solution to the **Subgraph Isomorphism Problem (SIP)** answers the question: is there a subgraph G' in G such that H can be mapped one-to-one, nodes and links, to G' ?

CDP is NP-hard (in 4 steps)

Given two graphs G and H , a solution to the **Subgraph Isomorphism Problem (SIP)** answers the question: is there a subgraph G' in G such that H can be mapped one-to-one, nodes and links, to G' ?

How to reduce this to CDP in poly-time?

CDP is NP-hard (in 4 steps)

Given two graphs G and H , a solution to the **Subgraph Isomorphism Problem (SIP)** answers the question: is there a subgraph G' in G such that H can be mapped one-to-one, nodes and links, to G' ?

Step 1 Change all vertices v of G into Fog nodes in I with available resources set to 1.

Step 2 Change all edges in (u, v) of G into node-to-node links in I with available bandwidth set to 1.

Step 3 Change all v of H into components of A with resource requirements set to 1.

Step 4 Change all edges (u, v) of H into component-component interactions in A with bandwidth requirement set to 1.

With Steps 1-4 we can do $(\text{SIP}) \rightarrow_p (\text{CDP})$. Hence $\text{CDP} \in \text{NP-hard}$.

* For all details of the proof see: A. Brogi, S.Forti, [QoS-aware Deployment of IoT Applications Through the Fog](#), 2017.

FogTorchII

1. solves CDP via **backtracking** search (exp-⌚, alas!)
2. estimates QoS-assurance by varying links QoS in 🎰 Monte Carlo simulations
3. estimates 📈 Fog resource consumption and 💸 monthly deployment cost



What about application deployment security?

Our contribution

A declarative methodology to assess the security level of multi-component application deployments in Fog scenarios.



Security Capabilities

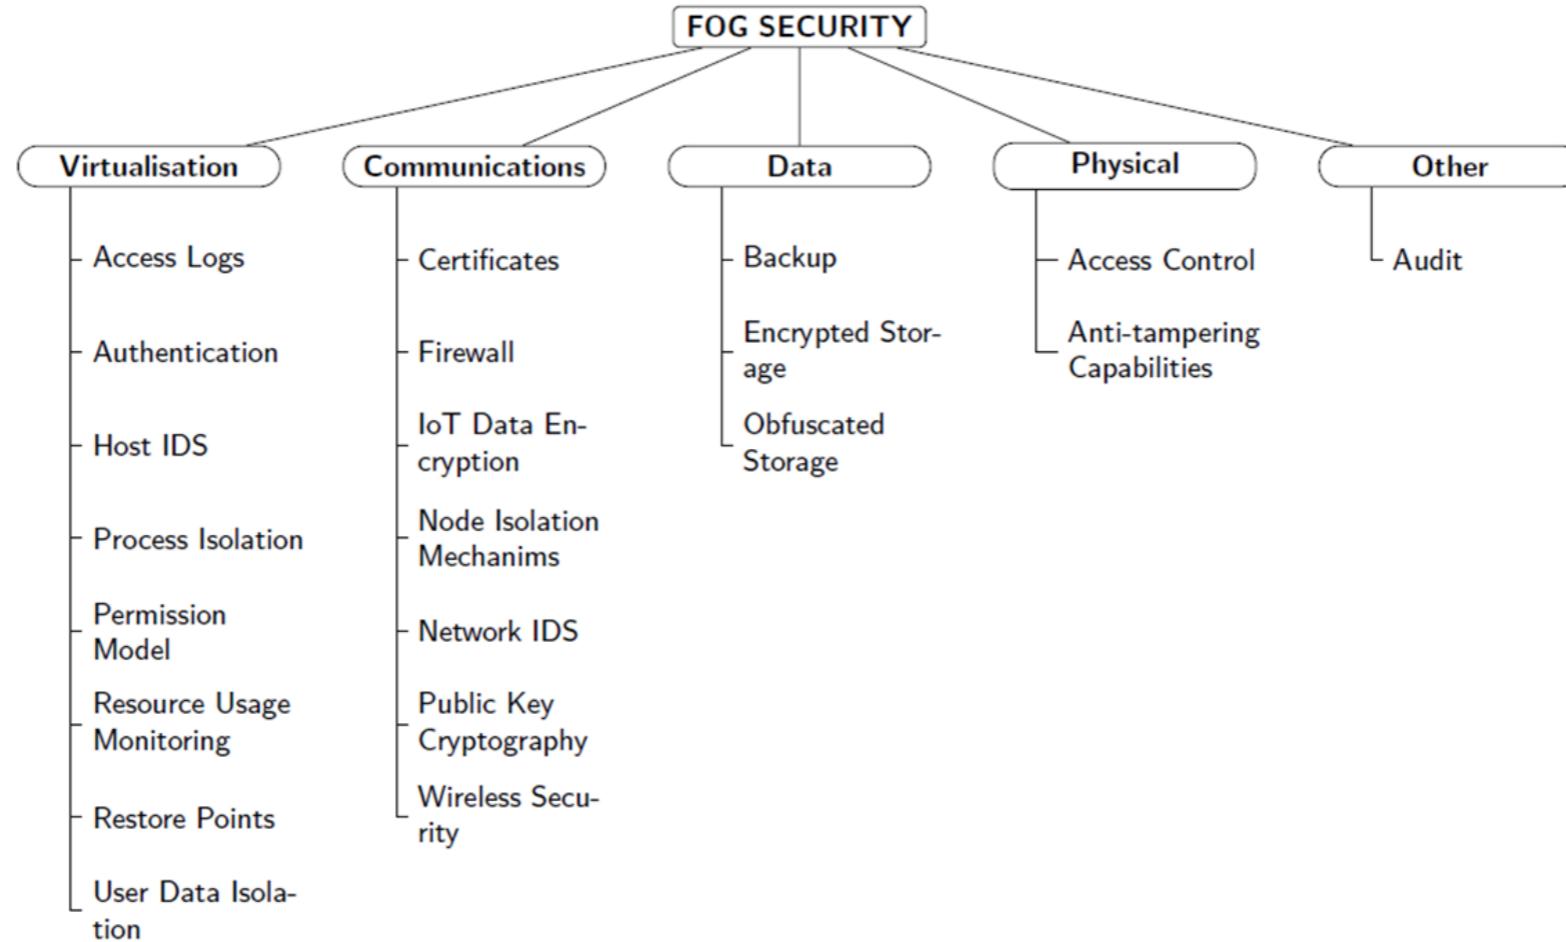
SecFog needs a **vocabulary** to specify **security capabilities** available/required in Fog scenarios.

Security control frameworks for the Cloud exists  (e.g., ISO/IEC 19086, EU Cloud SLA Standardisation Guidelines).

This is not the case for the Fog  

A Taxonomy

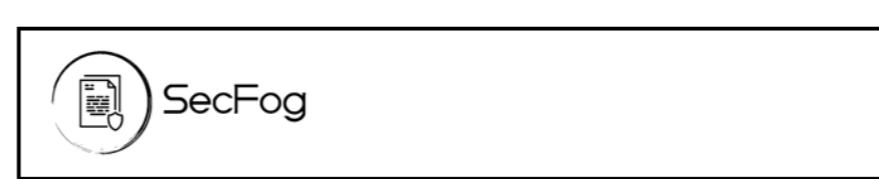
We hence proposed *our* taxonomy (based on recent surveys 



Big Picture

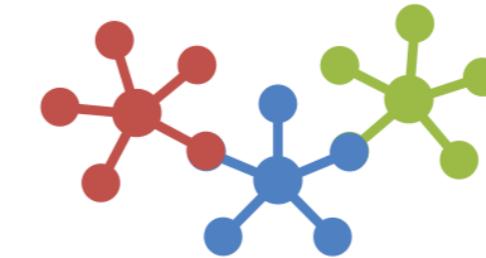
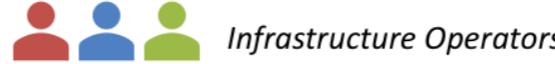


Infrastructure Operators



How does SecFog work?

Big Picture

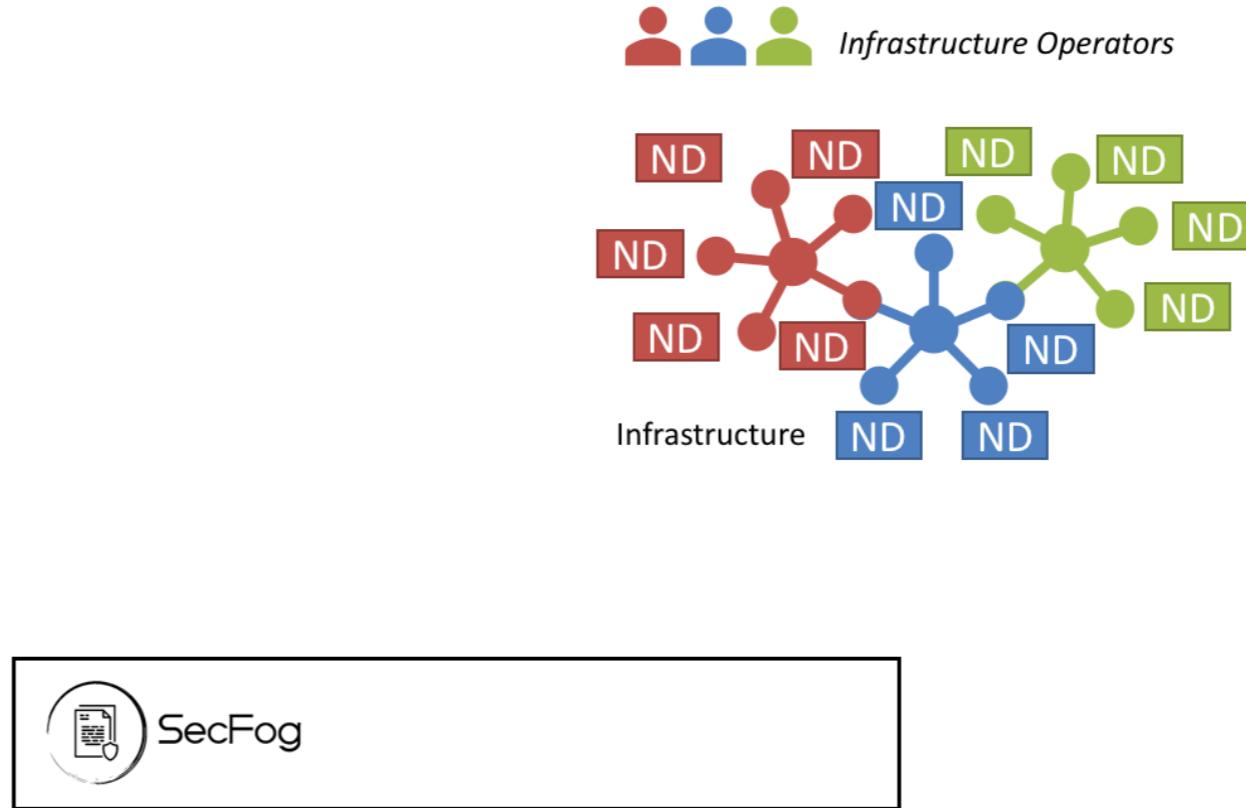


Infrastructure



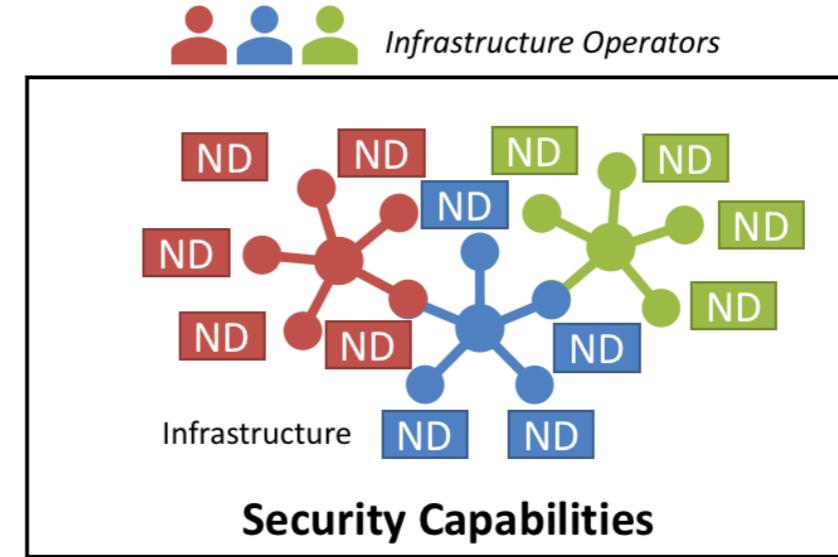
Different (small, medium, large) **operators** manage the Fog infrastructure.
Indeed, Fog deployments will span various service providers.

Big Picture



We assume that each node will self-describe its capabilities (and their *reliability*) through a **Node Descriptor** (ND) using the taxonomy vocabulary.

Big Picture



ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



In this way, we get a complete description of the available **Security Capabilities**.

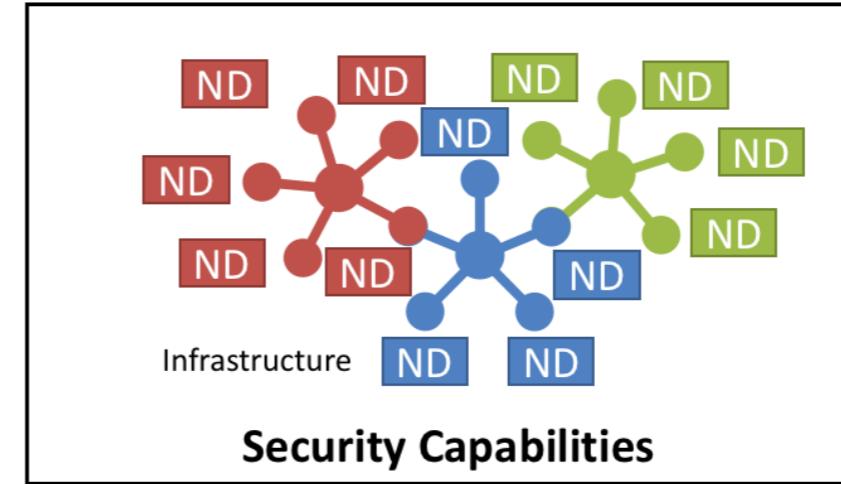
Big Picture



App Operator



Infrastructure Operators



ND = Node Descriptors

CR = Component Reqs

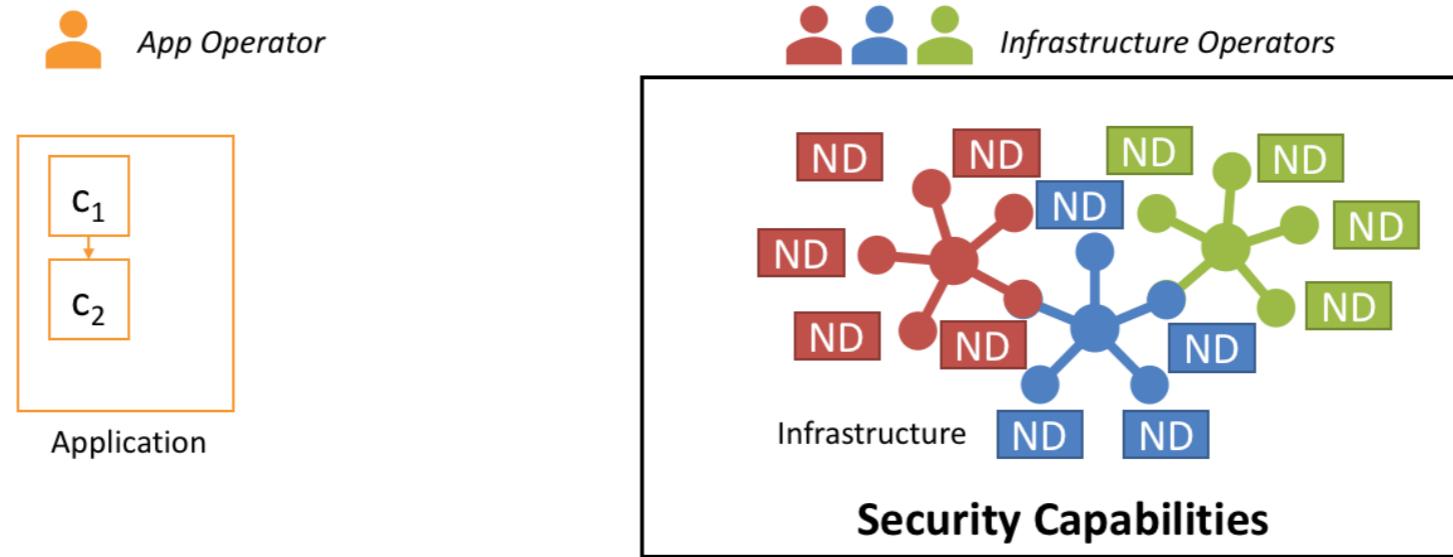
AR = Application Reqs



SecFog

On the other hand, we consider the **App Operator** that will describe their app.

Big Picture

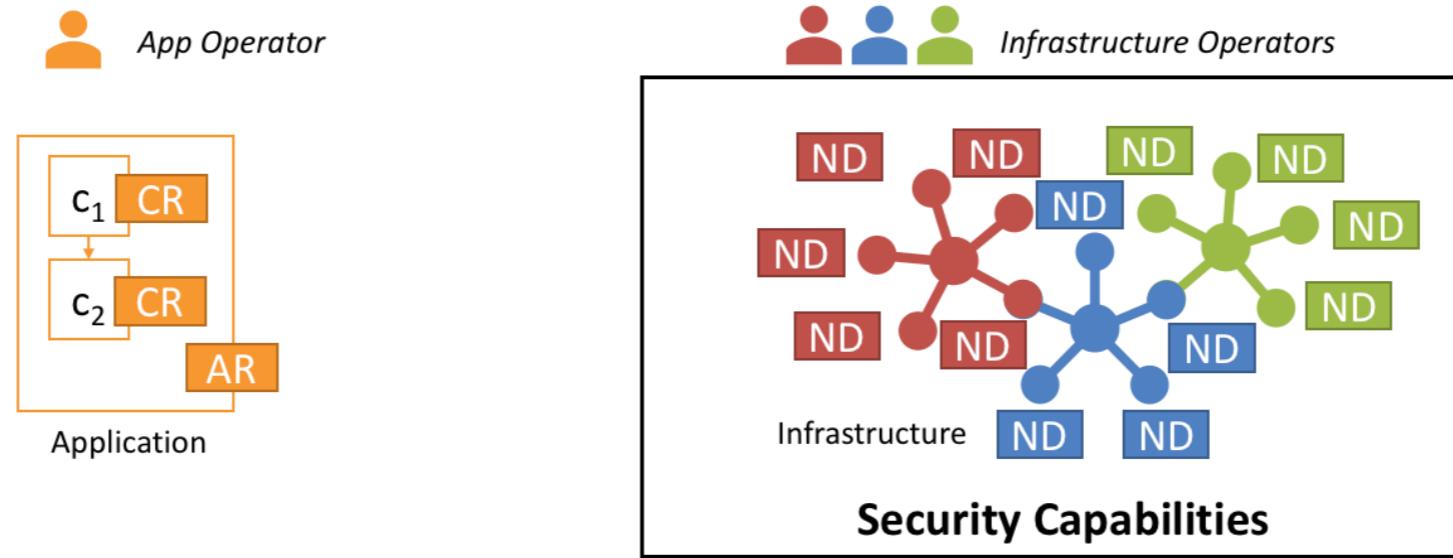


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



This can be done by specifying the **app topology**...

Big Picture

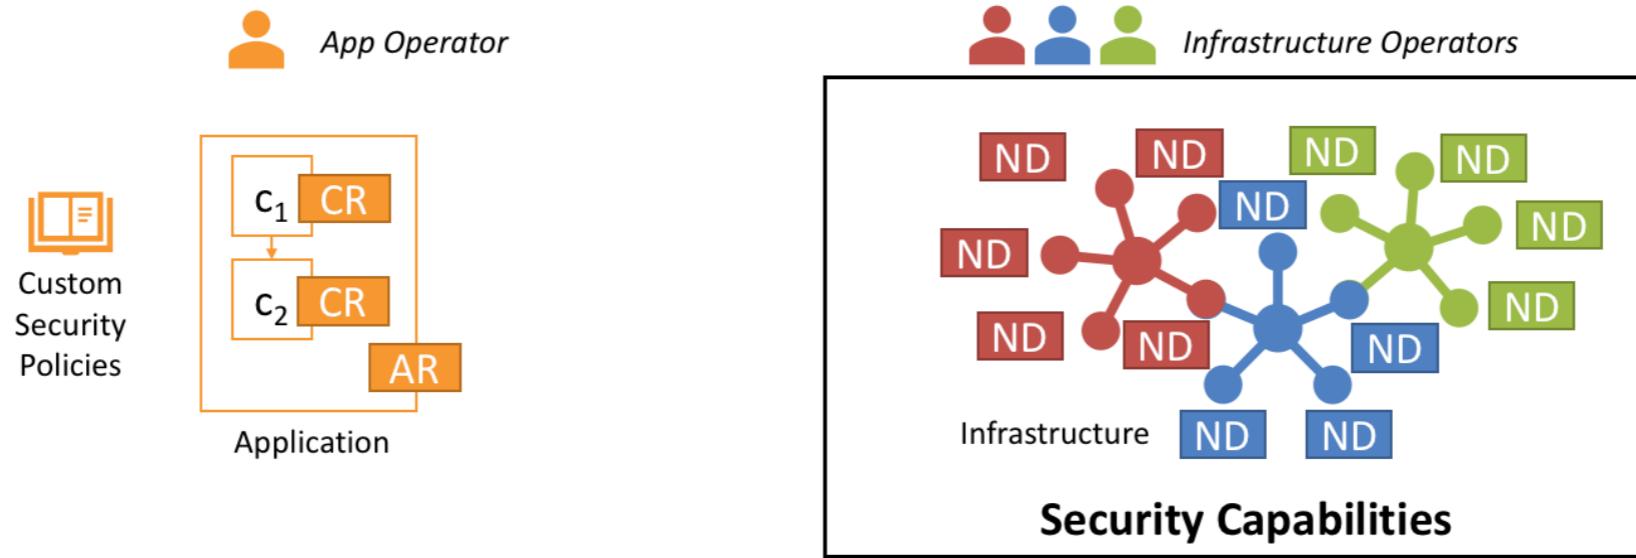


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



... and by annotating it with security requirements both at the **component (CR)** and **application (AR)** level, again relying on the taxonomy dictionary.

Big Picture

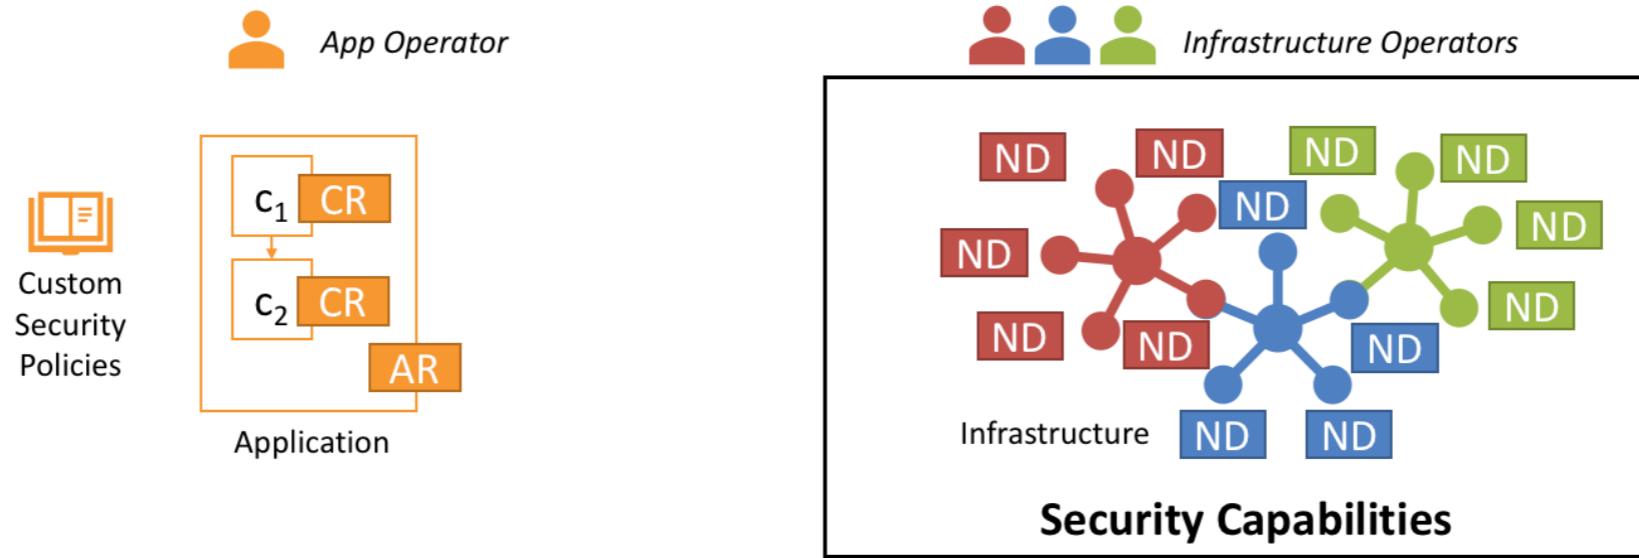


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



AR and CR can be expressed in term of **Custom Security Policies**.

Big Picture

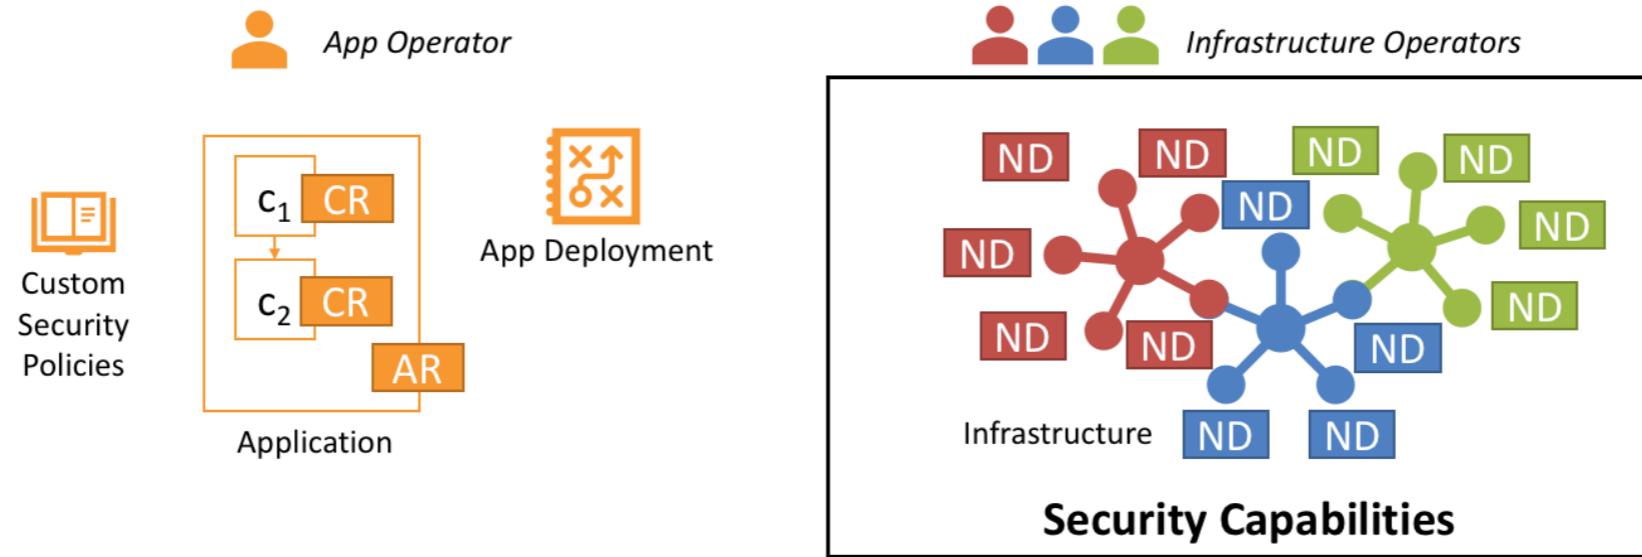


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



Custom Security Policies are expressed in terms of the **Default Security Policies** of SecFog.

Big Picture

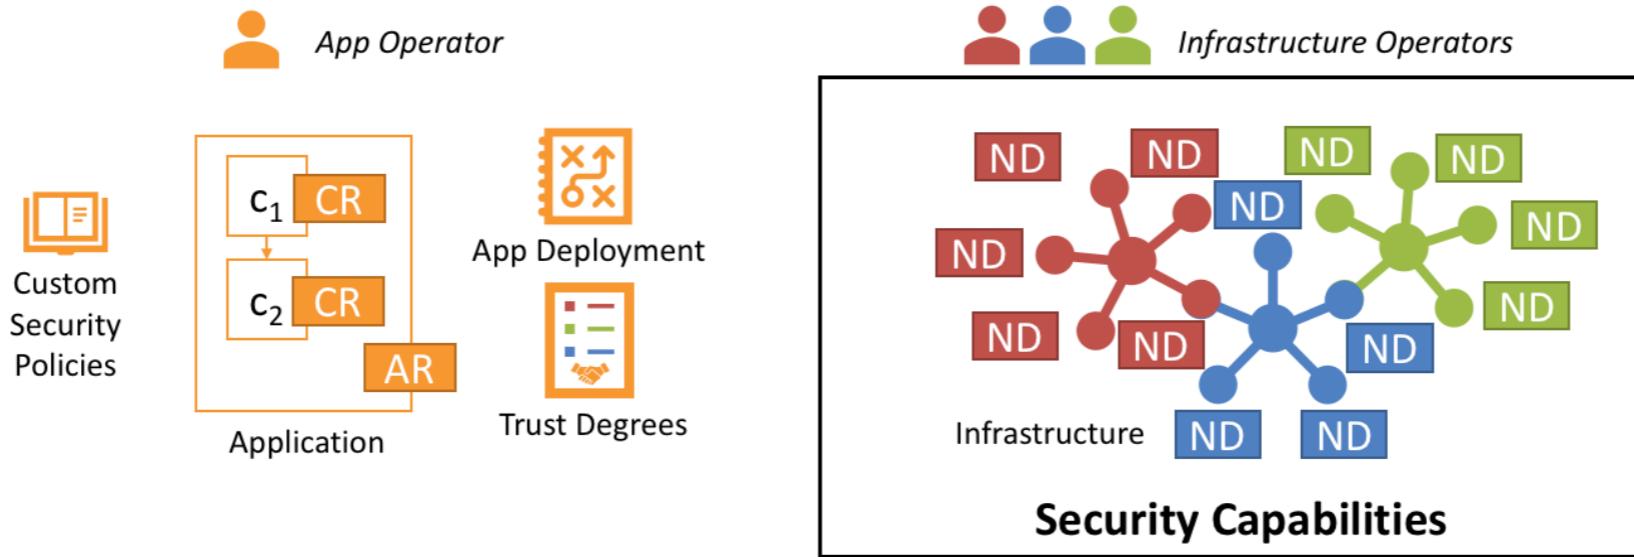


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



The App Operator can also specify complete or partial App Deployments of their application.

Big Picture

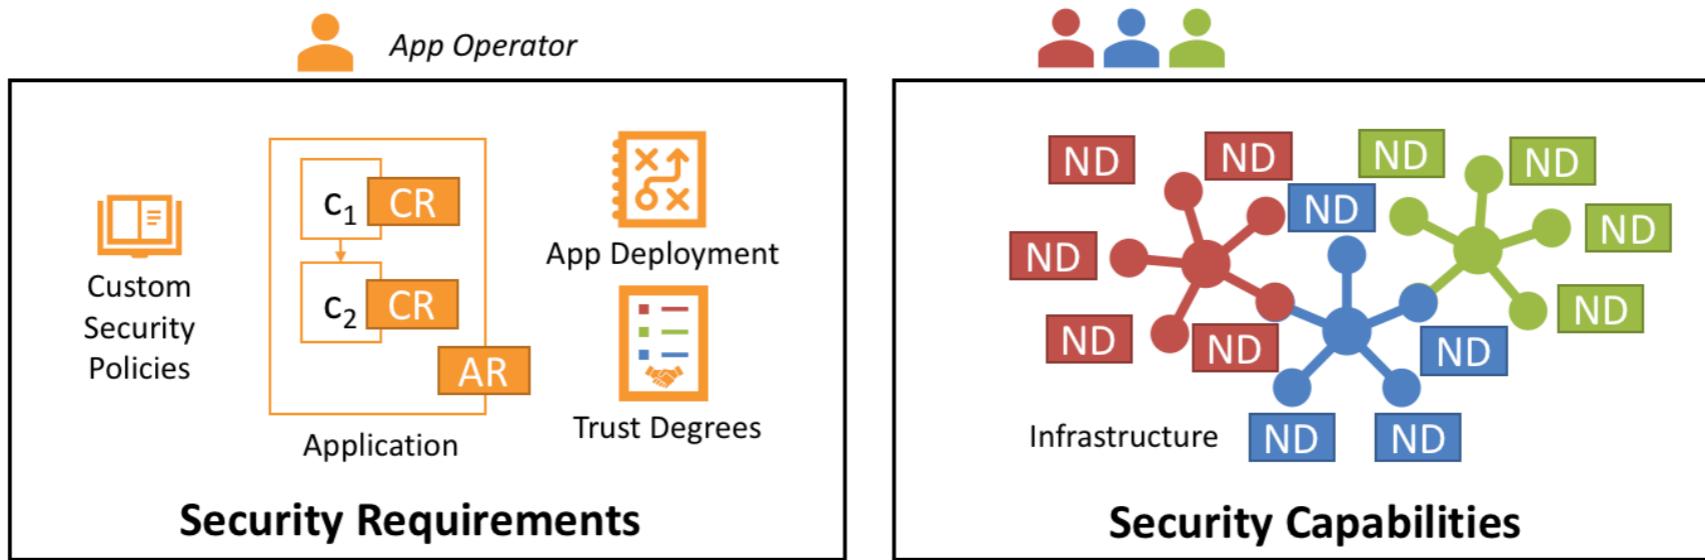


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



Finally, the App Operator can specify the **Trust Degrees** towards different Infrastructure Operators.

Big Picture

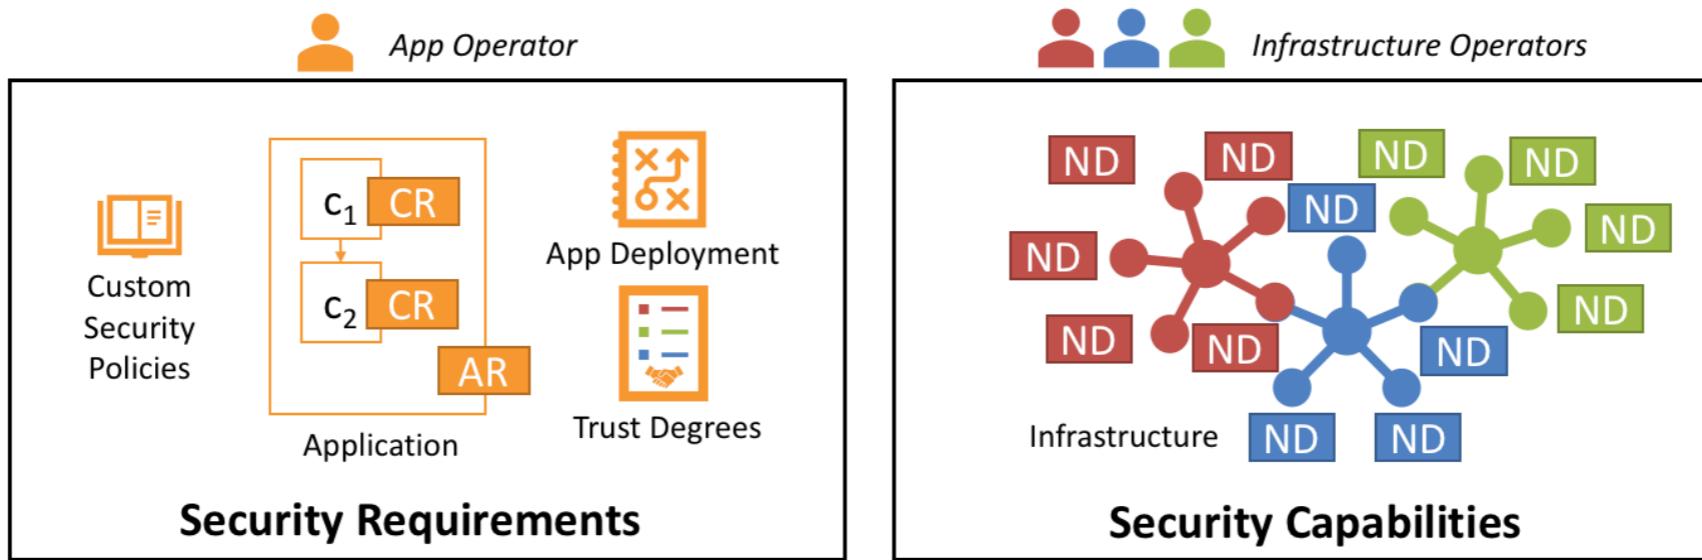


ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs



All this constitutes the **Security Requirements** of the considered multi-component app.

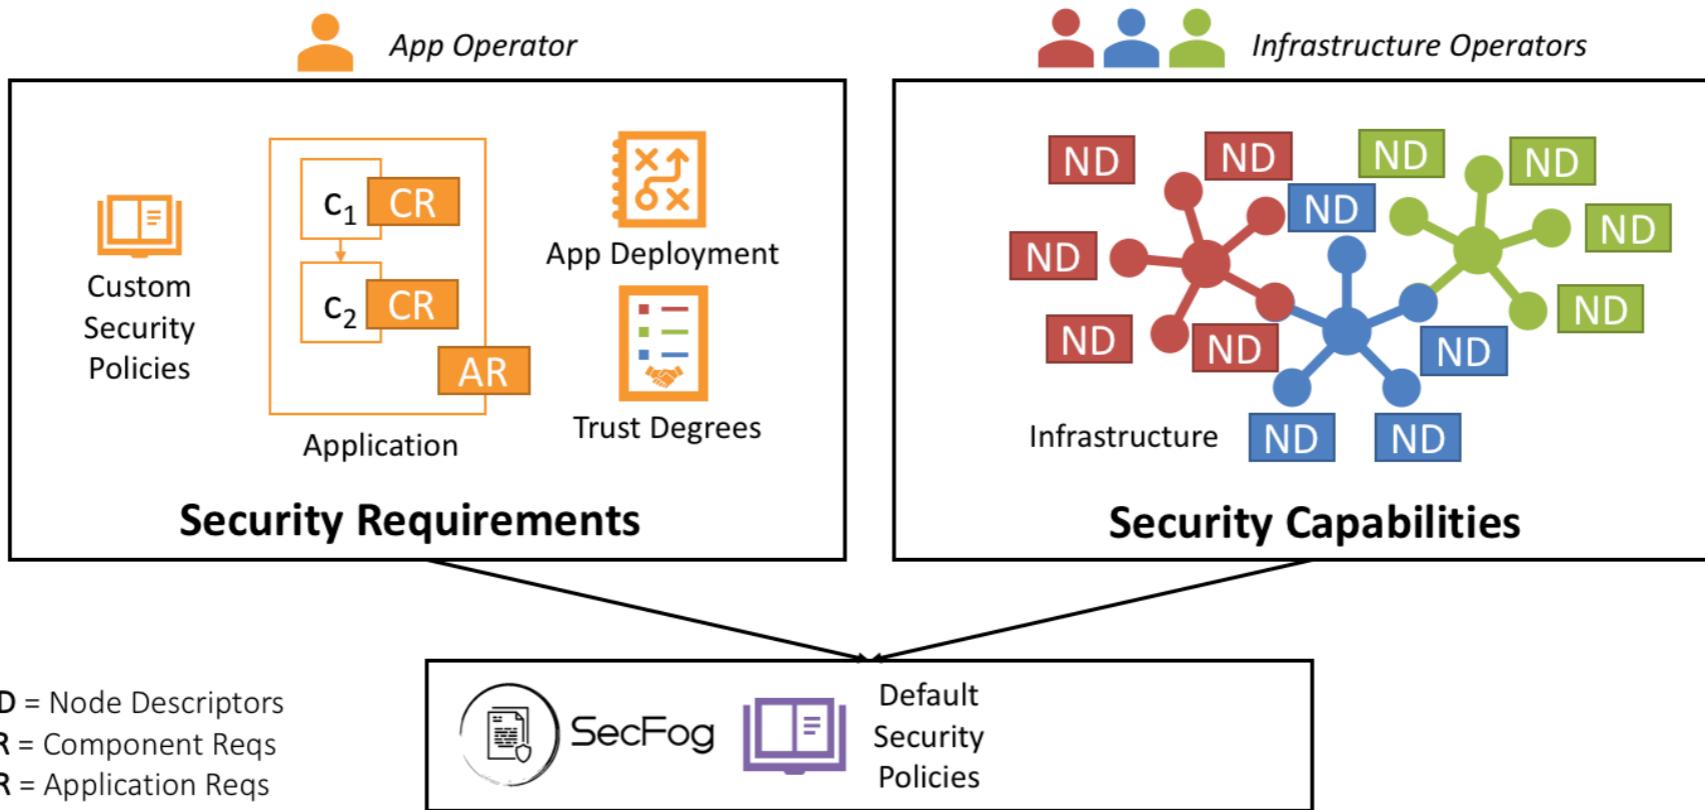
Big Picture



ND = Node Descriptors
CR = Component Reqs
AR = Application Reqs

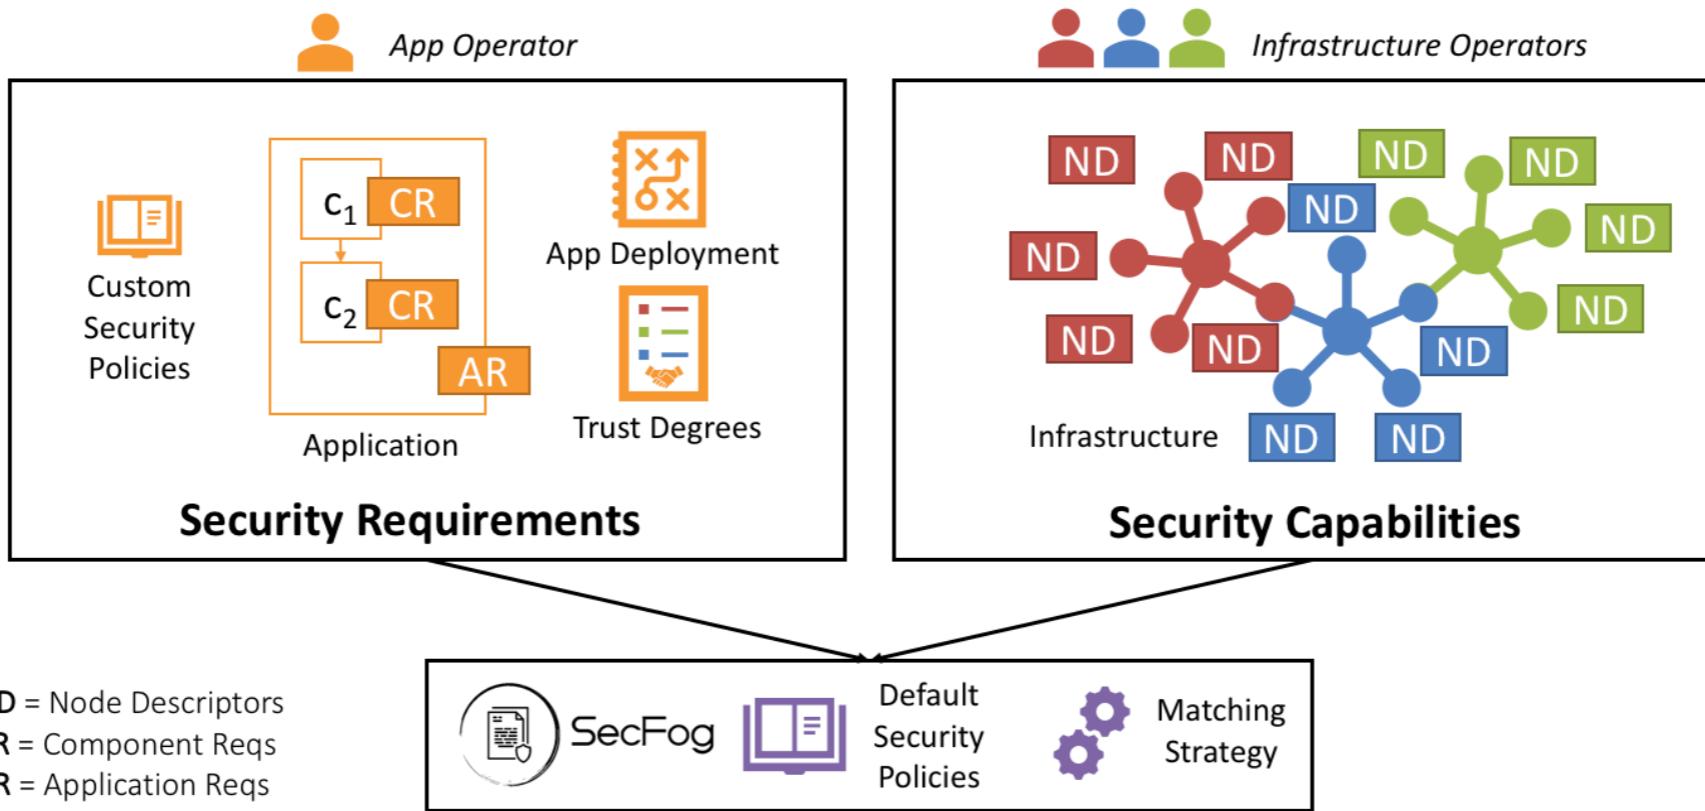
SecFog takes as input the Security Capabilities and...

Big Picture



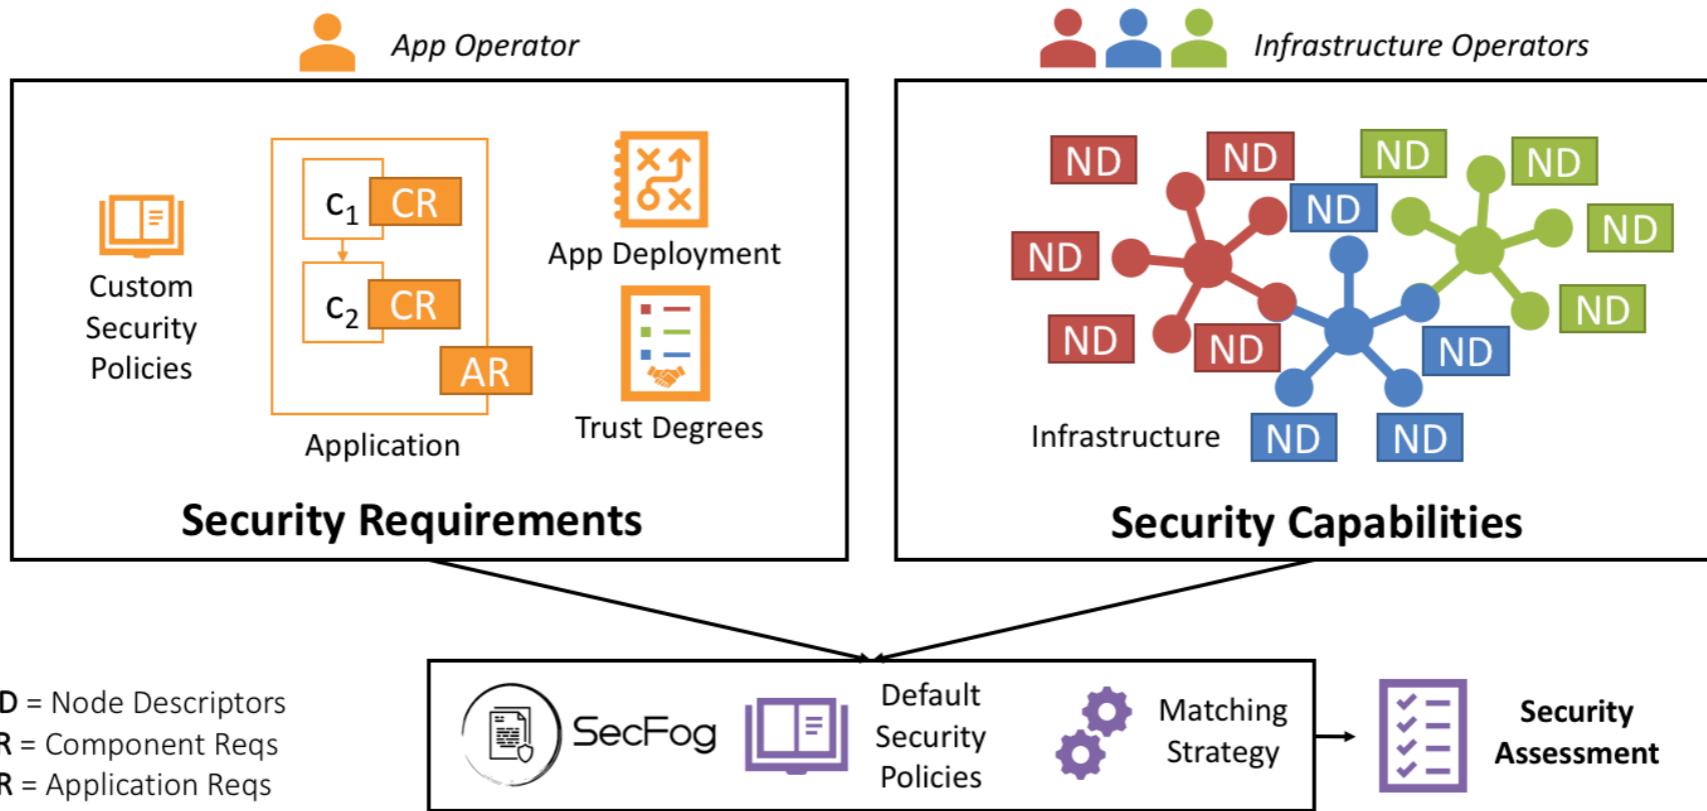
... the Security Requirements.

Big Picture



A matching strategy is used to determine secure deployments...

Big Picture



... and to assess their **Security Level**.

Implementation

SecFog has been prototyped using the **ProbLog2** language.

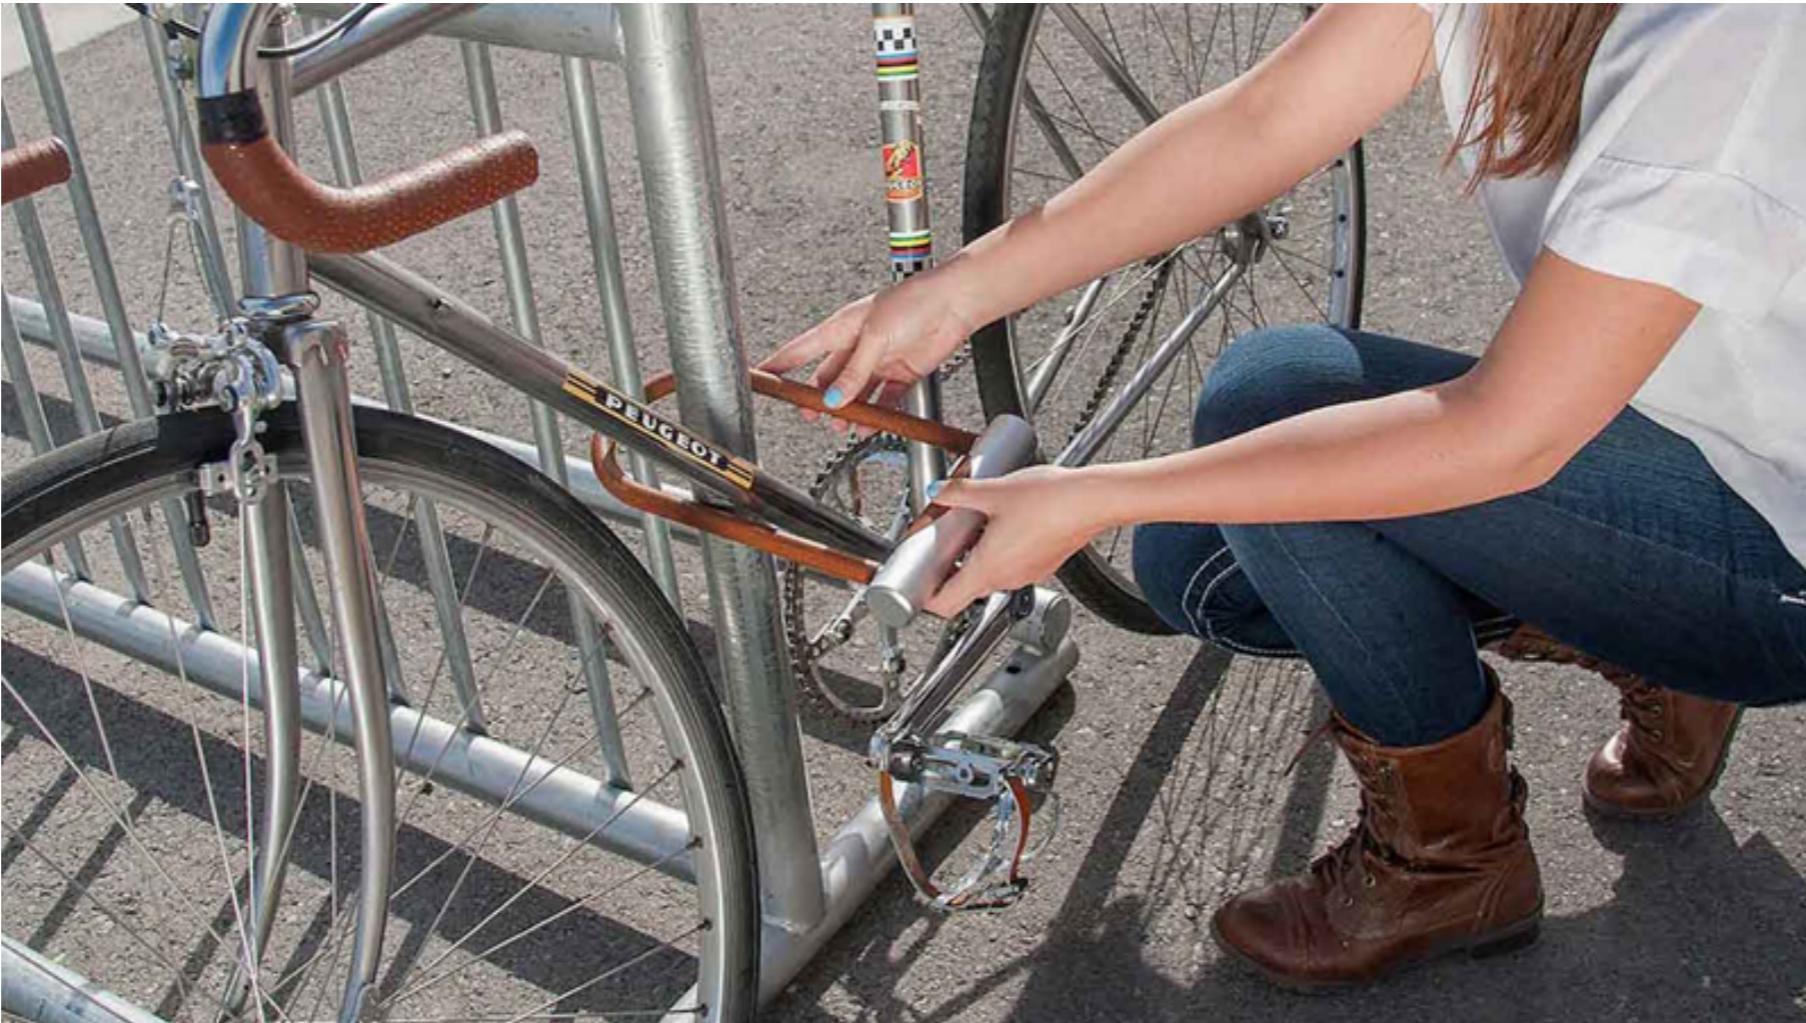
The prototype:

1. Matches application security requirements with infrastructure capabilities (*generate & test* approach), and
2. Obtains the security level of each deployment by multiplying the reliability of all exploited security capabilities, weighted by trust degrees.

It is **only** sixteen (16) lines of code 

( We are gonna write them all! )

A note on point (2) of the previous slide...



If I use a single lock for my bike, the security level of me locking my bike is the "certified" reliability of the lock \times my level of trust in the lock vendor (has (s)he got a cut of my key?)

A note on point (2) of the previous slide...



What if I'm in Pisa and I decide to use **more locks**?

Now change locks with security capabilities and vendor with infrastructure operator...

Problog 101

ProbLog is a Python package that permits writing logic programs that encode complex interactions between large sets of heterogeneous components, capturing the inherent uncertainties that are present in real-life situations.

Problog programs are composed of **facts***

```
p::f. % statement f is true with probability p
```

and **rules**

```
r :- c1, ..., cn. % property r is inferred when c1 & ... & cn hold
```

* A fact declared simply as `f.` is assumed to be `true` with probability `1`.

Problog 101

- ProbLog programs are logic programs in which some of the facts are annotated with (their) probabilities.
 - Each program defines a probability distribution over logic programs where a fact `p::f.` is considered `true` with probability `p` and false with probability `1-p`.
 - The ProbLog engine determines the success probability of a query `q` as the probability that `q` has a proof, given the distribution over logic programs.
- 🤓 ☕ We are ready to implement SecFog!

SecFog - Deployments

Given a list of app components $[c|cs]$, we can *recursively* define a **deployment** as the association of each component c to a node n , i.e. $d(c,n)$:

```
:- use_module(library(lists)).  
  
deployment([],[]). % an empty set of components Leads to an empty deployment  
deployment([C|Cs],[d(C,N)|D]) :-  
    node(N,_),  
    deployment(Cs,D).
```

We can try this (also [online](#)):

```
node(fog1, fogOp1). % node(id, operatorId)  
node(fog2, fogOp2).  
node(cloud1, cloudOp).  
  
query(deployment([iot_controller, data_storage, dashboard], _)).
```

SecFog - Secure Deployment

We first write a rule to check:

1. if a component `c` is deployed to an existing node `n` (as per deployment `d`), and
2. that the infrastructure operator `Op` managing `n` is trustable according to the application operator.

```
secure(C, N, D) :-  
    member(d(C,N), D), % is the couple d(C,N) in D?  
    node(N, Op),  
    trustable(Op).
```

Try it:

```
0.8::trustable(fogOp1).  
0.5::trustable(fogOp2).  
query(secure(iot_controller, _, [d(iot_controller, _)])).
```

SecFog - Secure Application Deployment

Now that we have a way to check that an app component is securely deployed, we write rules to recursively check that the application as a whole was securely deployed.

```
secureApp(A,D) :-  
    app(A, L),  
    deployment(L,D),  
    secureComponents(A,L,D).  
  
secureComponents(A, []).  
secureComponents(A, [C|Cs],D) :-  
    secureComponent(C,N,D),  
    secureComponents(A, Cs,D).
```

📦 That's all folks! These are the default security policies of SecFog.
As promised, they are only 16 lines of code (import's excluded 😊).

How to use SecFog

The app operator needs to define a `secureComponent(C,N,D)` for each app component, always including the predicate `secure(C, N, D)`.

Then, it is possible to run the query:

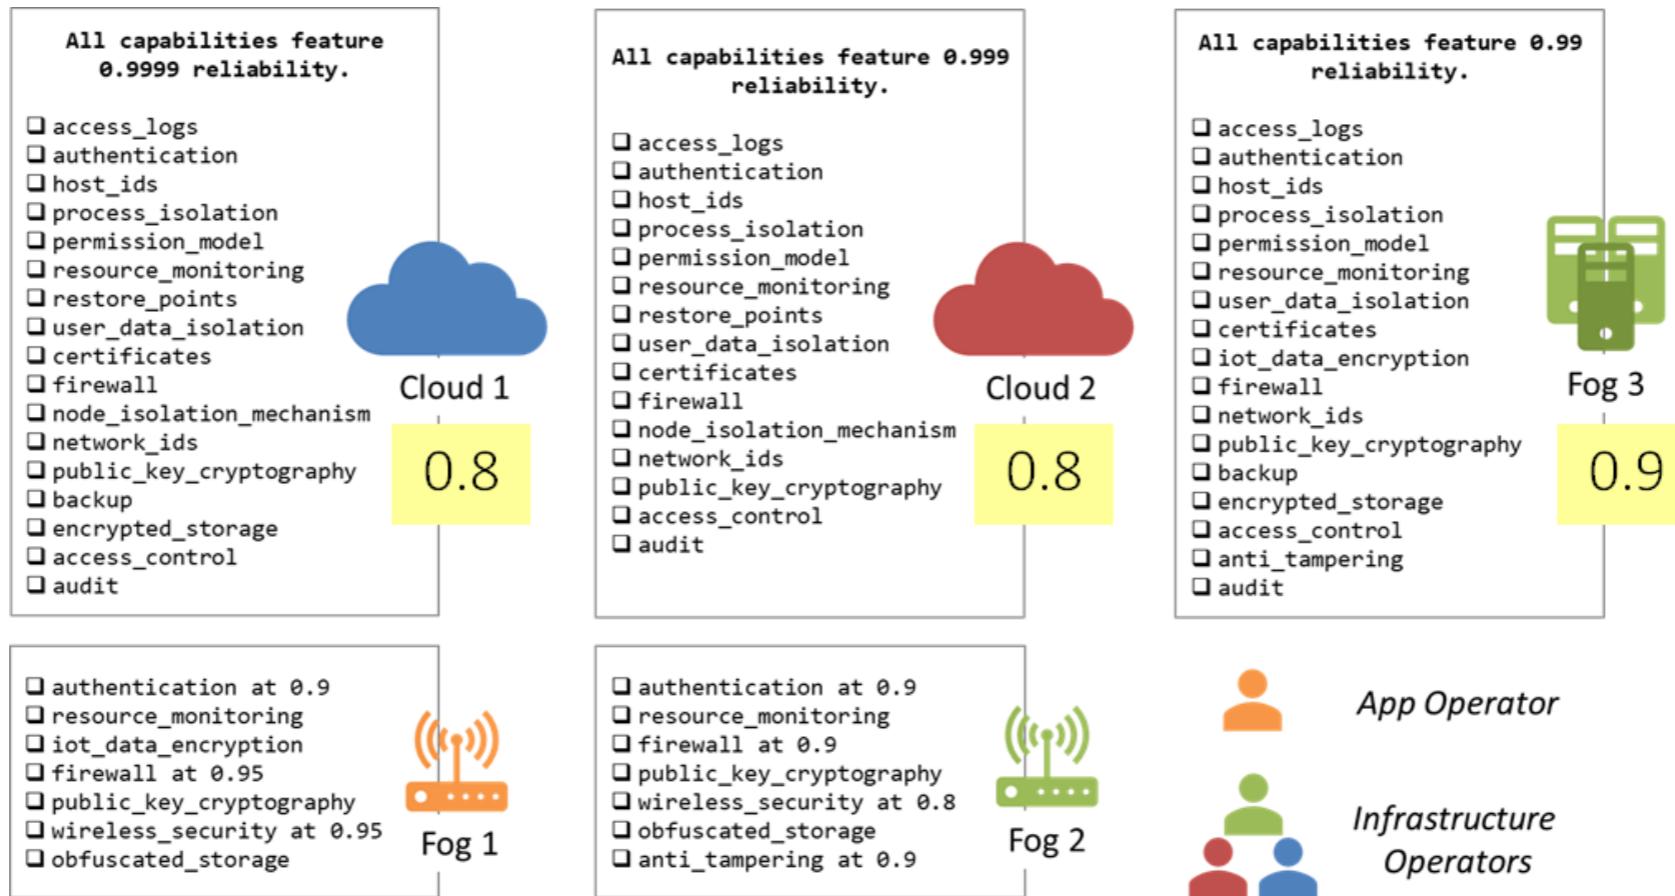
```
query(secureApp(A,_)).
```

to obtain all secure deployments and their security level.

Let's see a **complete example**.

Example: Node Descriptors

The file `infrastructure.pl` programmatically declares the Node Descriptors for the infrastructure skecthec below. How do we include trust degrees (as per the yellow boxes)?



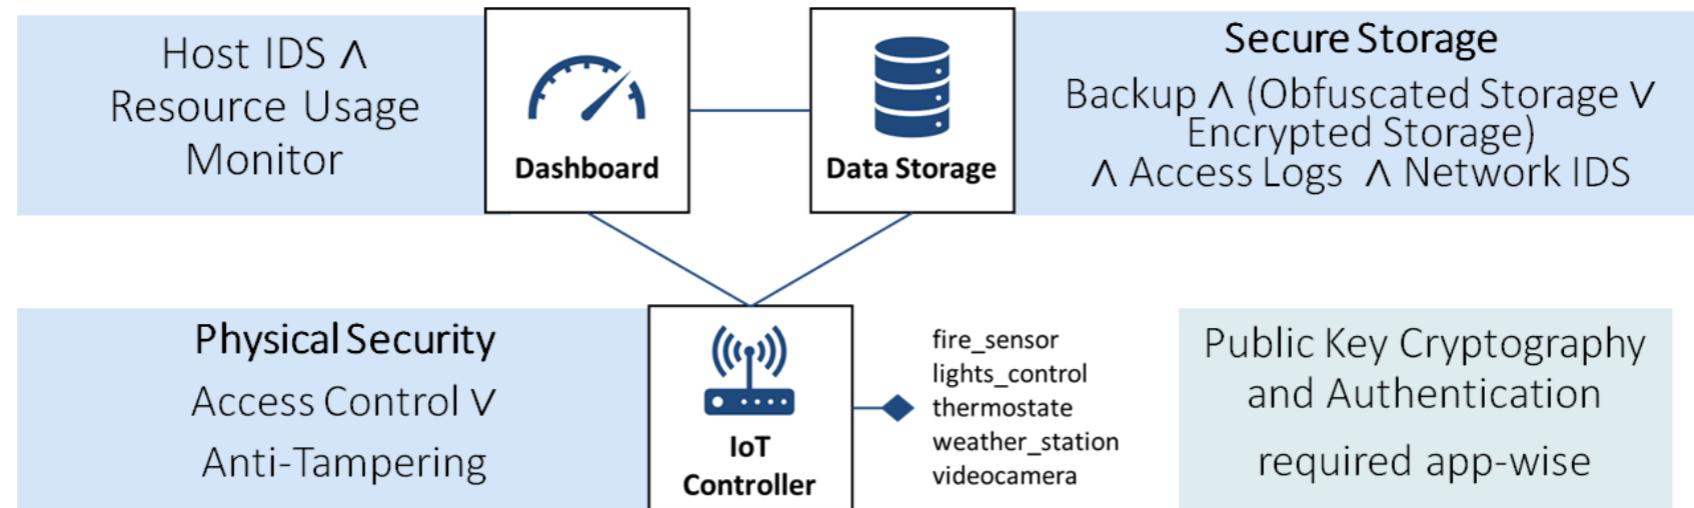
Example: Trust Degrees

Yes  as done before!

```
%trust degrees
0.8::trustable(cloudOp1).
0.8::trustable(cloudOp2).
0.9::trustable(fogOp).
trustable(appOp).
```

Example: Application

The application



can be specified as:

```
app(smartbuilding, [iot_controller, data_storage, dashboard]).
```

What about its **security requirements**?

Example: Component Security Requirements

The requirements of the IoT Controller component



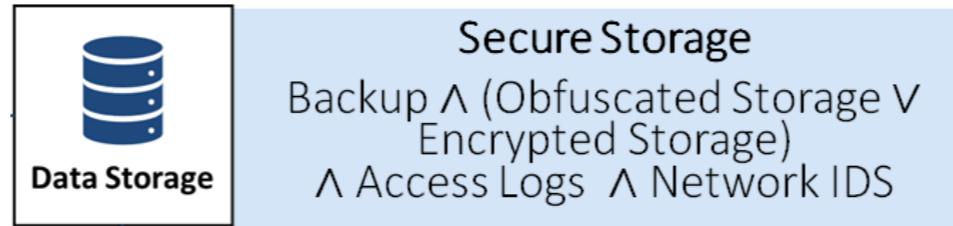
can be specified as:

```
physical_security(N) :- anti_tampering(N); access_control(N).
```

```
secureComponent(iot_controller, N, D) :-  
    physical_security(N),  
    secure(iot_controller, N, D).
```

Example: Component Security Requirements

The requirements of the DataStorage component



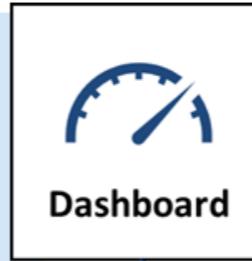
can be specified as:

```
secure_storage(N) :-  
    backup(N),  
    (encrypted_storage(N); obfuscated_storage(N)).  
  
secureComponent(data_storage, N, D) :-  
    secure_storage(N),  
    access_logs(N),  
    network_ids(N),  
    secure(data_storage, N, D).
```

? Pop Quiz ?

Can you write component security requirements for the Dashboard?

Host IDS ∧
Resource Usage
Monitor



Any volunteers?

```
secureComponent(dashboard, N, D) :-
```

Example: App Security Requirements

Finally, app-wise security requirements

Public Key Cryptography
and Authentication
required app-wise

can be specified as:

```
extras([]).  
extras([d(C,N)|Ds]) :-  
    public_key_cryptography(N),  
    authentication(N),  
    extras(Ds).
```

Example: Query

In this case, we (re-)define also:

```
mySecureApp(A,D) :-  
    secureApp(A,D),  
    deployment(L,D),  
    extras(D).
```

Which will be used to query the system on eligible deployments and security levels:

```
query(mySecureApp(A,_)).
```

(All `example.pl` code is now on Moodle)

To run the programmed specification, issue the command

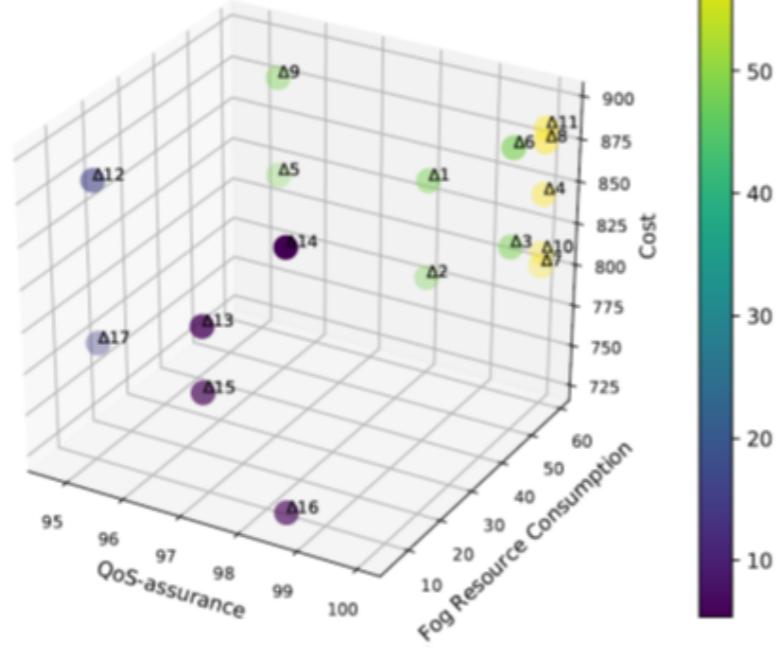
```
$ problog ./example.pl
```

SecFog Output



```
mySecureApp(smartbuilding,[d(iot_controller,cloud1), d(data_storage,cloud1), d(dashboard,cloud1)]): 0.79928029
mySecureApp(smartbuilding,[d(iot_controller,cloud1), d(data_storage,cloud1), d(dashboard,cloud2)]): 0.63699776
mySecureApp(smartbuilding,[d(iot_controller,cloud1), d(data_storage,cloud1), d(dashboard,fog3)]): 0.69114513
mySecureApp(smartbuilding,[d(iot_controller,cloud1), d(data_storage,fog3), d(dashboard,cloud1)]): 0.67752684
mySecureApp(smartbuilding,[d(iot_controller,cloud1), d(data_storage,fog3), d(dashboard,cloud2)]): 0.53996463
mySecureApp(smartbuilding,[d(iot_controller,cloud1), d(data_storage,fog3), d(dashboard,fog3)]): 0.66417689
mySecureApp(smartbuilding,[d(iot_controller,cloud2), d(data_storage,cloud1), d(dashboard,cloud1)]): 0.63757163
mySecureApp(smartbuilding,[d(iot_controller,cloud2), d(data_storage,cloud1), d(dashboard,cloud2)]): 0.63642441
mySecureApp(smartbuilding,[d(iot_controller,cloud2), d(data_storage,cloud1), d(dashboard,fog3)]): 0.55131415
mySecureApp(smartbuilding,[d(iot_controller,cloud2), d(data_storage,fog3), d(dashboard,cloud1)]): 0.54045108
mySecureApp(smartbuilding,[d(iot_controller,cloud2), d(data_storage,fog3), d(dashboard,cloud2)]): 0.67448315
mySecureApp(smartbuilding,[d(iot_controller,cloud2), d(data_storage,fog3), d(dashboard,fog3)]): 0.66238504
mySecureApp(smartbuilding,[d(iot_controller,fog2), d(data_storage,cloud1), d(dashboard,cloud1)]): 0.5827336
mySecureApp(smartbuilding,[d(iot_controller,fog2), d(data_storage,cloud1), d(dashboard,cloud2)]): 0.46441781
mySecureApp(smartbuilding,[d(iot_controller,fog2), d(data_storage,cloud1), d(dashboard,fog3)]): 0.55988355
mySecureApp(smartbuilding,[d(iot_controller,fog2), d(data_storage,fog3), d(dashboard,cloud1)]): 0.54885163
mySecureApp(smartbuilding,[d(iot_controller,fog2), d(data_storage,fog3), d(dashboard,cloud2)]): 0.54687823
mySecureApp(smartbuilding,[d(iot_controller,fog2), d(data_storage,fog3), d(dashboard,fog3)]): 0.67268088
mySecureApp(smartbuilding,[d(iot_controller,fog3), d(data_storage,cloud1), d(dashboard,cloud1)]): 0.70503715
mySecureApp(smartbuilding,[d(iot_controller,fog3), d(data_storage,cloud1), d(dashboard,cloud2)]): 0.56188935
mySecureApp(smartbuilding,[d(iot_controller,fog3), d(data_storage,cloud1), d(dashboard,fog3)]): 0.69114513
mySecureApp(smartbuilding,[d(iot_controller,fog3), d(data_storage,fog3), d(dashboard,cloud1)]): 0.67752684
mySecureApp(smartbuilding,[d(iot_controller,fog3), d(data_storage,fog3), d(dashboard,cloud2)]): 0.67509079
mySecureApp(smartbuilding,[d(iot_controller,fog3), d(data_storage,fog3), d(dashboard,fog3)]): 0.83038718
```

FogTorchII Output



Dep. ID	IoTController	DataStorage	Dashboard	QoS	Resources	Cost
Δ1	Fog 2	Fog 3	Cloud 2	98.6%	48.4%	€ 856.7
Δ2	Fog 2	Fog 3	Cloud 1	98.6%	48.4%	€ 798.7
Δ3	Fog 3	Fog 3	Cloud 1	100%	48.4%	€ 829.7
Δ4	Fog 2	Fog 3	Fog 1	100%	59.2%	€ 844.7
Δ5	Fog 1	Fog 3	Cloud 1	96%	48.4%	€ 837.7
Δ6	Fog 3	Fog 3	Cloud 2	100%	48.4%	€ 887.7
Δ7	Fog 3	Fog 3	Fog 2	100%	59.2%	€ 801.7
Δ8	Fog 3	Fog 3	Fog 1	100%	59.2%	€ 875.7
Δ9	Fog 1	Fog 3	Cloud 2	96%	48.4%	€ 895.7
Δ10	Fog 1	Fog 3	Fog 2	100%	59.2%	€ 809.7
Δ11	Fog 1	Fog 3	Fog 1	100%	59.2%	€ 883.7
Δ12*	Fog 2	Cloud 2	Fog 1	94.7%	16.1%	€ 870.7
Δ13*	Fog 2	Cloud 2	Cloud 1	97.2%	5.4%	€ 824.7
Δ14*	Fog 2	Cloud 2	Cloud 2	98.6%	5.4%	€ 882.7
Δ15*	Fog 2	Cloud 1	Cloud 2	97.2%	5.4%	€ 785.7
Δ16*	Fog 2	Cloud 1	Cloud 1	98.6%	5.4%	€ 727.7
Δ17*	Fog 2	Cloud 1	Fog 1	94.7%	16.1%	€ 773.7

How to combine QoS-assurance, resource usage, cost and security? Any ideas?

Multi-objective Optimisation

We can rank deployments according to

$$r(\Delta) = \sum_{m \in M} \omega_m \times \overline{m(\Delta)}$$

where

- M is the set of metrics to be optimised,
- ω_m is the weight assigned to each metric, and
- $\overline{m(\Delta)}$ is the normalised value of the metric m for deployment Δ .

The better the deployment, the higher the ranking 

Results

Most probably, we aim at:

- *maximising* QoS-assurance and security
- *minimising* cost

For Fog resource usage we might go either:

- **Cloud-ward** = minimise Fog resource usage
- **Fog-ward** = maximise Fog resource usage

Dep. ID	IoTController	DataStorage	Dashboard	$r_F(\Delta)$	$r_C(\Delta)$
$\Delta 1$	Fog 2	Fog 3	Cloud 2	0.53	0.28
$\Delta 2$	Fog 2	Fog 3	Cloud 1	0.63	0.38
$\Delta 3$	Fog 3	Fog 3	Cloud 1	0.85	0.60
$\Delta 6$	Fog 3	Fog 3	Cloud 2	0.75	0.50
$\Delta 15^*$	Fog 2	Cloud 1	Cloud 2	0.15	0.40
$\Delta 16^*$	Fog 2	Cloud 1	Cloud 1	0.51	0.76

Conclusions



SecFog

help(s)

Fog app
operators who

want to determine reducing
secure app manual
deployments by tuning and

**considering app reqs,
infrastructure
capabilities and trust**

unlike existing
approaches*

* With which SecFog can be synergically used, as shown with FogTorchΠ.

Future Work

-  Better quantify trust degrees and reliability (perhaps using a semi-ring other than probability)
-  Improve scalability over large infrastructures via meta-heuristics (e.g., bio-inspired).
-  Engineer the whole FogTorchII + SecFog methodology into a (micro-)service and experiment with them in lifelike settings.



Yes, if you liked the ~~Fog~~ course... you can ask us for a thesis!

An (Interactive) Introduction to



SecFog

Stefano Forti

Department of Computer Science, University of Pisa, Italy

Reading: A. Brogi, S. Forti, G.-L. Ferrari, [Secure Apps in the Fog: Anything to Declare?](#), 2018.

This slideset was realised using Marp: <https://github.com/yhatt/marp>